

NIST AI 600-1 Applicability Determinations

Sentinel Nexus — Non-Applicable GAI Risk Families

RW

June 2026

Contents

Purpose	2
Profile-role basis	2
Determinations	2
Re-evaluation triggers	3

Purpose

NIST AI 600-1 expects an organization to **record the determination** even when a GAI risk is judged out of scope (GV-1.3-003 calls for a written test/response policy). This document states, per risk family, why several of the profile's *misuse / content-generation* risks are **not applicable** to Sentinel Nexus, so an auditor sees a reasoned decision rather than an omission.

Profile-role basis

Sentinel Nexus is a defensive SOC platform: it **consumes adversary telemetry and emits internal SOC verdicts and SOAR actions**. It is *not* a public-facing content generator — it does not produce text, images, audio, or video for human consumption or distribution. This posture is the common basis for the determinations below.

Determinations

NIST §2 risk	Determination	Rationale
2.1 CBRN Information or Capabilities	Not applicable	The system does not answer open-domain user questions and is not a chemical/biological design tool. Models are fine-tuned for SOC reasoning over fixed telemetry schemas; they do not synthesize CBRN knowledge. No general-purpose chat surface is exposed to users.
2.3 Dangerous, Violent, or Hateful Content	Not applicable	No content-generation surface. Outputs are structured verdicts (contain/monitor/dismiss) and forensic incident reports for internal responders, constrained by VerdictSchema/SoarExecutionSchema.

NIST §2 risk	Determination	Rationale
2.11 Obscene, Degrading, or Abusive Content (CSAM/NCII)	Not applicable	No image/audio/video generation; no public content output. Training corpora are synthetic SOC tradecraft + sanitized telemetry, curated and credential-scrubbed; no possibility of CSAM/NCII generation.
2.10 Intellectual Property	Largely not applicable	The system does not reproduce or generate copyrighted creative works. Residual consideration: frontier-model API terms and training-data provenance — addressed by the data-curation policy and the AI System Inventory.
2.8 Information Integrity (content provenance / synthetic media)	Partially applicable / mostly N-A	The profile's content-provenance / deepfake-detection concerns target public synthetic media, which Nexus does not produce. The applicable slice — provenance/traceability of the swarm's own reasoning — is covered by AI-origin-stamped incident reports and the audit ledgers.

Re-evaluation triggers

These determinations are **re-opened** if the platform's scope changes such that any of the following becomes true (MP-4.1-008 — re-evaluate on new-domain use):

- A general-purpose, user-facing generative/chat surface is exposed.
- The system begins producing content for human consumption or external distribution.
- Models are adapted to a new domain outside SOC telemetry reasoning.
- Multimodal (image/audio/video) generation is introduced.

Until such a trigger occurs, organizational risk-management effort is concentrated on the **applicable** high-exposure families — Confabulation, Information Security, Data

Privacy, Harmful Bias & Homogenization, Human-AI Configuration, and Value Chain
— per the NIST AI 600-1 control tracker.